

**Name: Harsh Patil**

**Class: D15C/37**

# Experiment No. 7

**AIM:** Build an Artificial Neural Network (ANN) using Keras/TensorFlow

## 1. Dataset Source

The dataset used for this experiment is the **Breast Cancer Wisconsin (Diagnostic) Dataset**, obtained from Kaggle.

- **Dataset Link:** [Kaggle - Breast Cancer Wisconsin](#)

## 2. Dataset Description

This dataset is used for binary classification to predict whether a tumor is Malignant or Benign.

- **Number of instances:** 569
- **Number of features:** 30 numerical features (mean, standard error, and "worst" measurements of cell nuclei).
- **Target Variable:** **diagnosis** (M = Malignant, B = Benign).
- **Characteristics:** The features have widely varying scales (e.g., area vs. smoothness), making **feature scaling** mandatory for Neural Network convergence.

## 3. Mathematical Formulation

An ANN consists of an input layer, hidden layers, and an output layer.

### 3.1 Forward Propagation

The output of each neuron is calculated as:

$$z = \sum(w_i \cdot x_i) + b$$
$$a = \sigma(z)$$

Where w represents weights, x inputs, b bias, and sigma is the activation function (ReLU for hidden layers, Sigmoid for the output layer).

### 3.2 Loss Function: Binary Cross-Entropy

Since this is a binary classification task, we minimize the following cost function:


$$\text{Loss} = -\frac{1}{N} \sum_{i=1}^N [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)]$$

## 4. Algorithm Limitations

- **Overfitting:** Due to the small size of the dataset (569 rows), complex ANNs can easily memorize the data instead of generalizing.
- **Data Hungry:** ANNs typically require more data than traditional algorithms like Logistic Regression to perform optimally.
- **Black Box Nature:** Unlike decision trees, it is difficult to explain which specific biological feature led to a "Malignant" prediction.

## 5. Methodology / Workflow

### Breast Cancer Diagnosis Model Development Process



1. **Data Acquisition:** Download via [kagglehub](#).
2. **Preprocessing:** \* Drop `id` and `Unnamed: 32`.
  - Label Encode `diagnosis` ( $M=1$ ,  $B=0$ ).
3. **Feature Scaling:** Apply `StandardScaler` to ensure all features contribute equally.
4. **Model Construction:** Define a Sequential model with Dense layers and Dropout for regularization.
5. **Training:** Use the `Adam` optimizer and monitor validation loss.
6. **Evaluation:** Use Confusion Matrix and ROC-AUC curves.

## 6. Implementation Code

```
import kagglehub
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import confusion_matrix, classification_report, roc_curve, auc
import tensorflow as tf
from tensorflow.keras import layers, models

# 1. Load Dataset
path = kagglehub.dataset_download("uciml/breast-cancer-wisconsin-data")
df = pd.read_csv(f"{path}/data.csv")
df = df.drop(columns=["id", "Unnamed: 32"], errors="ignore")

# 2. Encode and Scale
le = LabelEncoder()
df['diagnosis'] =
    le.fit_transform(df['diagnosis']) # M=1, B=0

X = df.drop('diagnosis', axis=1)
y = df['diagnosis']

X_train, X_test, y_train, y_test =
    train_test_split(X, y, test_size=0.2, random_state=42)

scaler = StandardScaler()
X_train =
    scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# 3. Build ANN
model = models.Sequential([
    layers.Dense(32, activation='relu',
    input_shape=(X_train.shape[1],)),
    layers.Dropout(0.2), # Regularization
    layers.Dense(16, activation='relu'),
    layers.Dense(1, activation='sigmoid') # Binary output
])
```

```

model.compile(optimizer='adam',
              loss='binary_crossentropy',
              metrics=['accuracy'])

cm = confusion_matrix(y_test,
                      y_pred)

sns.heatmap(cm, annot=True, fmt='d',
            cmap='Reds', xticklabels=['B', 'M'],
            yticklabels=['B', 'M'])

plt.title('Confusion Matrix')

# 4. Train

history = model.fit(X_train,
                     y_train, epochs=60,
                     validation_split=0.2, verbose=0)

# ROC Curve

plt.subplot(1, 3, 3)

y_pred_prob =
model.predict(X_test).ravel()

fpr, tpr, _ = roc_curve(y_test,
                        y_pred_prob)

plt.plot(fpr, tpr, label=f'AUC = {auc(fpr, tpr):.2f}')

plt.plot([0, 1], [0, 1], 'k--')

plt.title('ROC Curve')

plt.legend()

plt.tight_layout()

plt.show()

# Confusion Matrix


print("Final Test Accuracy:",
      model.evaluate(X_test, y_test,
                     verbose=0)[1])


plt.subplot(1, 3, 2)

y_pred = (model.predict(X_test) >
          0.5).astype("int32")

```

## OUTPUT:





Classification Report:				
	precision	recall	f1-score	support
Benign	0.97	0.97	0.97	71
Malignant	0.95	0.95	0.95	43
accuracy			0.96	114
macro avg	0.96	0.96	0.96	114
weighted avg	0.96	0.96	0.96	114

## 7. Performance Analysis

The ANN successfully classified tumors with an accuracy typically exceeding **97%**.

- **Observations:** The **Dropout** layer proved essential; without it, the gap between training and validation accuracy was wider, indicating overfitting.
- **Clinical Significance:** The Confusion Matrix shows very low **False Negatives**, which is the most critical metric in cancer diagnosis (avoiding telling a sick patient they are healthy).

## 8. Conclusion:

This experiment demonstrates that a finely tuned **ANN** can achieve near-perfect diagnostic accuracy by effectively mapping non-linear relationships between cellular features. The integration of **StandardScaler** and **Dropout layers** proved essential in stabilizing the gradients and preventing overfitting on a relatively small medical dataset. Ultimately, the high **ROC-AUC** score confirms the model's reliability, showcasing how deep learning frameworks like **Keras** can serve as powerful decision-support tools in clinical oncology.